



NANYANG  
TECHNOLOGICAL  
UNIVERSITY  
SINGAPORE

# CZ3005

## Artificial Intelligence

### Reinforcement Learning

Assoc Prof Bo AN

*Research area:* artificial intelligence,  
computational game theory, reinforcement  
learning, optimization


[www.ntu.edu.sg/home/boan](http://www.ntu.edu.sg/home/boan)

*Email:* [boan@ntu.edu.sg](mailto:boan@ntu.edu.sg)


*Office:* N4-02b-55



# Lesson Outline



- Some RL algorithms:
  - Monte-Carlo
  - Temporal difference
  - Q-learning
  - Deep Q-Network





# Reinforcement Learning

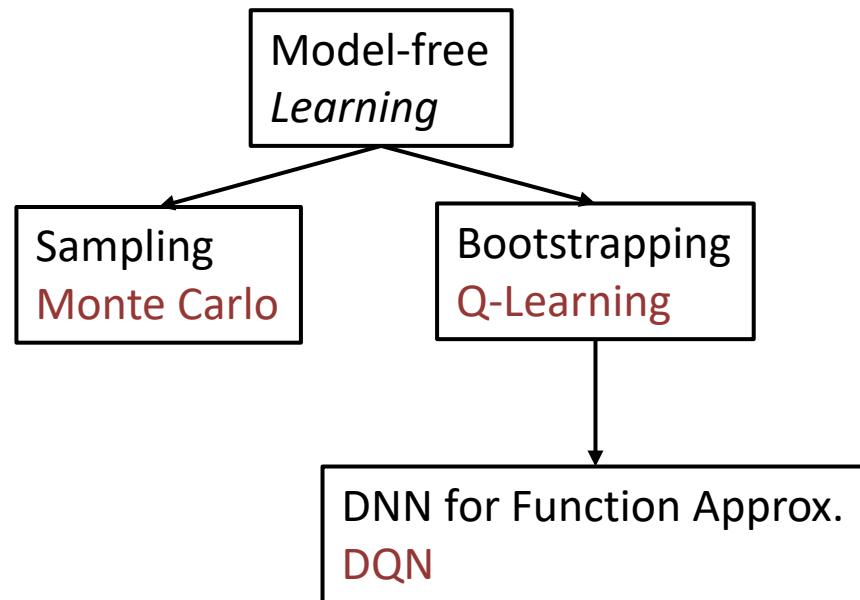
- Motivation
  - In last lecture, we compute **the value function** and find **the optimal policy**
  - But if without the transition function  $P(s'|s, a)$ ?
  - We can learn **the value function** and find **the optimal policy** without transition
    - From experience





# RL algorithms


- Types
  - Monte Carlo
  - Q-Learning
  - DQN
  - ...





# What is Monte Carlo

- Idea behind MC:
  - Just use randomness to solve a problem
- Simple definition:
  - Solve a problem by generating suitable random numbers and observing the fraction of numbers obeying some properties
- An example for calculating  $\pi$  (not policy in RL):
  - $S_{red} = \frac{1}{4}\pi r^2, S_{square} = r^2$
  - putting dots on the square randomly for  $n = 3000$  times
  - $\pi \approx 4 \times \frac{N_{red}}{n}$ ,  $N_{red}$  is the number of dots in the circle






# Monte Carlo in RL: Prediction

- Basic Idea: we run in the world randomly and gain experience to learn
- What experience? Many trajectories!
  - $(s_1, a_1, r_2, s_2, a_2, r_3, \dots, s_T), \dots$
- What we learn? Value function!
  - Recall that the return is the total discounted rewards:
$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^n r_{t+n} + \dots = \sum_i \gamma^i r_{t+i}$$
  - Recall that the value function is the expected return from  $s$ 
$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$
- How we learn?
  - Use experience to learn an empirical state value function  $\tilde{V}_{\pi}(s) = \frac{1}{N} \sum_{i=1}^N G_{i,s}$



# An Example

- One-dimensional grid world
  - A robot is in a 1x4 world
  - State: current cell  $s \in [cell_1, cell_2, cell_3, cell_4]$
  - Action: left or right
  - Reward:
    - Move one step (-1)
    - Reach the destination cell (+10) (ignoring the one-step reward)



Start point




Destination

# One-dimensional Grid World



- Trajectory or episode:

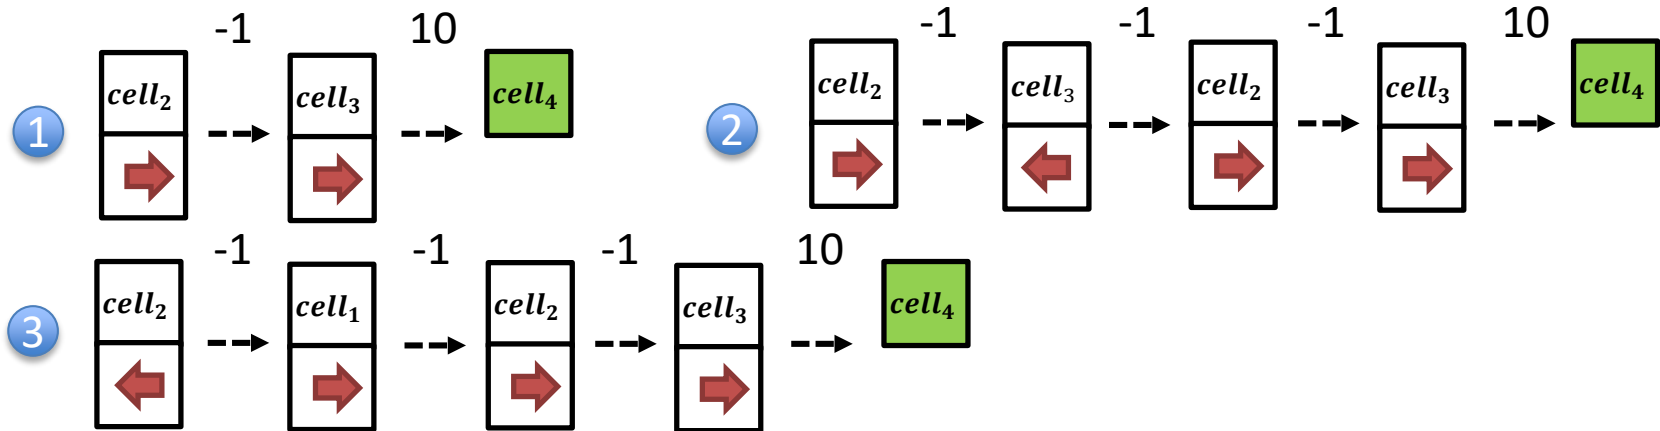
- The sequence of states from the starting state to the terminal state
- Robot starts in  $cell_2$ , ends in  $cell_4$



Start point

Destination

- The representation of the three episodes








# Compute Value Function

- Idea: Average return observed after visits to  $(s, a)$
- First-visit MC: average returns only for **first** time  $(s, a)$  is visited in an episode
- Return in one episode (trajectory):

$$G_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots + \gamma^n r_{t+n} + \dots = \sum_i \gamma^i r_{t+i}$$

- We calculate the return for  $cell_2$  of first episode with  $\gamma = 0.9$




$$G_t = -1 \times 0.9^0 + 10 \times 0.9^1 = 9$$




# Compute Value Function (cont'd)

- Similarly the return for  $cell_2$  of second episode with  $\gamma = 0.9$



$$G_t = -1 \times 0.9^0 - 1 \times 0.9^1 - 1 \times 0.9^2 + 10 \times 0.9^3 = 4.58$$

- Similarly the return for  $cell_2$  of third episode with  $\gamma = 0.9$



$$G_t = -1 \times 0.9^0 - 1 \times 0.9^1 - 1 \times 0.9^2 + 10 \times 0.9^3 = 4.58$$

- The empirical value function for  $cell_2$  is  $\frac{9 + 4.58 + 4.58}{3} = 6.0533 \dots$



# Compute Value Function (cont'd)

- Given these three episodes, we compute the value function for all non-terminal state

6.2	6.05	8.73
$cell_1$	$cell_2$	$cell_3$

- We can get more accurate value function with more episodes

# First Visit Monte Carlo Policy Evaluation

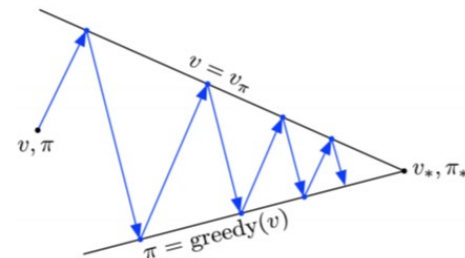
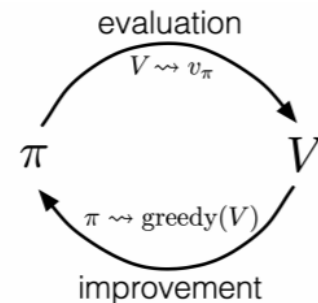


- Average returns only for the first time  $s$  is visited in an episode
- Algorithm
  - Initialize:
    - $\pi \leftarrow$  policy to be evaluated
    - $V \leftarrow$  an arbitrary state-value function
    - $Returns(s) \leftarrow$  an empty list, for all state  $s$
  - Repeat many times:
    - Generate an episode using  $\pi$
    - For each state  $s$  appearing in the episode:
      - $R \leftarrow$  return following **the first occurrence** of  $s$
      - Append  $R$  to  $Returns(s)$
      - $V(s) \leftarrow average(Returns(s))$



# Monte Carlo in RL: Control

- Now, we have the value function of all states given a policy
- We need to improve policy to be better
- Policy Iteration
  - Policy evaluation
  - Policy improvement
- However, we need to know how good an action is






# Q-value

- Estimate how good an action is when staying in a state
- Defined as the expected return starting from  $s$ , taking the action  $a$  and thereafter following policy  $\pi$

$$Q^\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$$


- Representation: A table
  - Filled with the Q-value given a state and an action





# Computing Q-value


- MC for estimating Q:
  - A slight difference from estimating the value function
  - Average returns for state-action pair  $(s, a)$  is visited in an episode
- We calculate the return for  $(cell_2, \text{right})$  of first episode with  $\gamma = 0.9$






# Compute Q-Value (cont'd)

- Similarly the return for ( $cell_2$ , right) of second episode with  $\gamma = 0.9$



- Similarly the return for ( $cell_2$ , right) of third episode with  $\gamma = 0.9$



- The empirical Q-value function for ( $cell_2$ , right) is  $\frac{9 + 4.58 + 9}{3} = 7.53$













# Q-Value for Control

- Filling the Q-table
  - By going through all state-action pairs, we get a complete Q-table with all the entries filled
  - A possible Q-table example

*cell<sub>1</sub> cell<sub>2</sub> cell<sub>3</sub>*

	8.1	9.3	9.9
	4.5	5.6	7.5
			
			

- Selecting action

$$\pi'(s) = \operatorname{argmax}_{a \in A} Q^\pi(s, a)$$

At *cell<sub>1</sub>*, *cell<sub>2</sub>* and *cell<sub>3</sub>*, we choose right



# MC control algorithm

Initialize, for all  $s \in \mathcal{S}$ ,  $a \in \mathcal{A}(s)$ :

$Q(s, a) \leftarrow$  arbitrary

$Returns(s, a) \leftarrow$  empty list

$\pi \leftarrow$  an arbitrary  $\varepsilon$ -soft policy

Repeat forever:

(a) Generate an episode using  $\pi$

(b) For each pair  $s, a$  appearing in the episode:

$R \leftarrow$  return following the first occurrence of  $s, a$

Append  $R$  to  $Returns(s, a)$

$Q(s, a) \leftarrow \text{average}(Returns(s, a))$

(c) For each  $s$  in the episode:

$a^* \leftarrow \arg \max_a Q(s, a)$

For all  $a \in \mathcal{A}(s)$ :

$$\pi(s, a) \leftarrow \begin{cases} 1 - \varepsilon + \varepsilon / |\mathcal{A}(s)| & \text{if } a = a^* \\ \varepsilon / |\mathcal{A}(s)| & \text{if } a \neq a^* \end{cases}$$

Policy evaluation

Policy improvement



# Q-Learning

- Previously, we need the whole trajectory
- In Q-Learning, we only need one-step trajectory:  $(s, a, r, s')$
- The difference is the Q-value computing

– Previously:

$$\tilde{Q}_{\pi}(s, a) = \frac{1}{N} \sum_{i=1}^N G_{i,s}$$

– Now, updating rule:

$$Q_{new}(S_t, A_t) \leftarrow Q_{old}(S_t, A_t) + \alpha (R_{t+1} + \gamma \max_a Q_{old}(S_{t+1}, a) - Q_{old}(S_t, A_t))$$





# Q-Learning

Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$



$S \leftarrow S'$

    until  $S$  is terminal



# A Step-by-step Example


- 5-room environment as MDP
  - We'll number each room 0 through 4
  - The outside of the building can be thought of as one big room 5
  - End at room 5
  - Notice that doors at rooms 1 and 4 lead into the building from room 5 (outside)





# A Step-by-step Example (cont'd)

- Goal
  - Put an agent in any room, and from that room, go outside (or room 5)
- Reward
  - The doors that lead immediately to the goal have an instant reward of 100
  - Other doors not directly connected to the target room have zero reward



$$R = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} & \text{action} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 100 \\ 0 & 0 & 0 & 0 & 0 & 100 \end{bmatrix} \end{matrix}$$


state



# Q-Learning Step by Step

- Initialize matrix Q as a zero matrix
- $\alpha = 0.01, \gamma = 0.99$
- Loop for each episode until converge
  - Initial state: current we are in room 1 (1<sup>st</sup> outer loop)
  - Loop for each step of episode (until reach room 5)
    - ... (Next slide)

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$






# Q-Learning Step by Step (cont'd)

- ... (last slide)
  - Loop for each step of episode (until room 5)
    - By random selection, we go to 5
    - We get 100 reward
    - Update Q:  $Q_{new}(S_t, A_t) \leftarrow Q_{old}(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q_{old}(S_{t+1}, a) - Q_{old}(S_t, A_t))$ 
      - At room 5, we have 3 possible actions: go to 1, 4 or 5; We select the one with max reward
      - $Q_{new}(1,5) \leftarrow Q_{old}(1,5) + \alpha \left( 100 + \gamma \max_a Q_{old}(5, a) - Q_{old}(1,5) \right) = 0 + 0.01 \times (100 + 0.99 \times 0 - 0) = 1$

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$



$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & \mathbf{1} \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix} \end{matrix}$$






# Q-Learning Step by Step (cont'd)

- When we loop many episodes, we can get

$$Q = \begin{matrix} & \begin{matrix} 0 & 1 & 2 & 3 & 4 & 5 \end{matrix} \\ \begin{matrix} 0 \\ 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 0 & 80 & 0 \\ 0 & 0 & 0 & 64 & 0 & 100 \\ 0 & 0 & 0 & 64 & 0 & 0 \\ 0 & 80 & 51 & 0 & 80 & 0 \\ 64 & 0 & 0 & 64 & 0 & 100 \\ 0 & 80 & 0 & 0 & 80 & 100 \end{bmatrix} \end{matrix}$$

- According to this Q-table, we can select actions
  - E.g. We are at room 2
  - Greedily select based on maximum of Q value





# An Example of Iteration Process

- A complex grid world example
- [https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld\\_td.html](https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html)



# Deep Q-Network

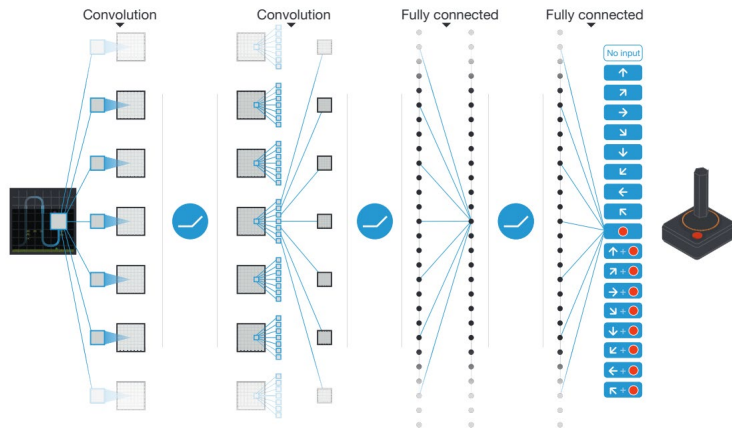

- Previously, we represent the Q-value as a table
- However, tabular representation is insufficient
  - Many real world problems have enormous state and/or action spaces
  - Backgammon:  $10^{20}$  states
  - Computer Go:  $10^{170}$  states
  - Robots: continuous state space
- We use a neural network as a black box to replace the table
  - Input a state and an action, output the Q-value






# DQN in Atari

- Input state  $s$  is stack of raw pixels from last 4 frames
- Output is  $q(s,a)$  for 18 button
- Reward is change in score for that step



# DQN in Atari (cont'd)



- Pong's video
- <https://www.youtube.com/watch?v=PSQt5KGv7Vk>
- Beat human on many games

